

Normal model with variance known

Suppose that we have a sample x_1, \dots, x_n and we assume the following model

$$\left\{ N(\mu, \sigma^2); \mu \in \mathbb{R} \right\}$$

That is a normal distribution with mean μ , unknown and variance σ^2 , known. We want to test

$$H_0 : \mu = \mu_0 \qquad \text{against} \qquad H_1 : \mu \neq \mu_0$$

The idea is to set up a test of hypothesis based on the Maximum Likelihood Estimator $\hat{\mu}$.

Likelihood Ratio Test (LRT)

By definition, the Likelihood Ratio Test (LRT) statistic is given by

$$\Lambda(x_1, \dots, x_n) = \frac{\prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{x_i - \hat{\mu}}{\sigma}\right)^2}}{\prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{x_i - \mu_0}{\sigma}\right)^2}}$$

Which simplifies to

$$\Lambda(x_1, \dots, x_n) = \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \hat{\mu})^2 + \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu_0)^2\right)$$

Where $\hat{\mu}$ is the Maximum Likelihood estimator for μ and is equal to \bar{x} .

Now, it is interesting to note that twice the logarithm of $\Lambda(x_1, \dots, x_n)$ gives us a convenient expression (see next slide)

Likelihood Ratio Test (LRT)

$$\begin{aligned}2\ln\left(\Lambda(x_1, \dots, x_n)\right) &= -\frac{1}{\sigma^2} \sum_{i=1}^n \left(x_i - \bar{x}\right)^2 + \frac{1}{\sigma^2} \sum_{i=1}^n \left(x_i^2 - \mu_0\right)^2 \\&= -\frac{1}{\sigma^2} \sum_{i=1}^n \left(x_i^2 - 2x_i\bar{x} + \bar{x}^2\right) + \frac{1}{\sigma^2} \sum_{i=1}^n \left(x_i^2 - 2x_i\mu_0 + \mu_0^2\right) \\&= -\frac{1}{\sigma^2} \left(\sum_{i=1}^n x_i^2 - 2\bar{x} \sum_{i=1}^n x_i + \sum_{i=1}^n \bar{x}^2\right) + \frac{1}{\sigma^2} \left(\sum_{i=1}^n x_i^2 - 2\mu_0 \sum_{i=1}^n x_i + \sum_{i=1}^n \mu_0^2\right) \\&= \frac{1}{\sigma^2} \left(-\sum_{i=1}^n x_i^2 + 2\bar{x} \sum_{i=1}^n x_i - \sum_{i=1}^n \bar{x}^2 + \sum_{i=1}^n x_i^2 - 2\mu_0 \sum_{i=1}^n x_i + \sum_{i=1}^n \mu_0^2\right) \\&= \frac{n}{\sigma^2} \left(-\frac{1}{n} \sum_{i=1}^n x_i^2 + 2\bar{x} \frac{1}{n} \sum_{i=1}^n x_i - \frac{1}{n} n\bar{x}^2 + \frac{1}{n} \sum_{i=1}^n x_i^2 - 2\mu_0 \frac{1}{n} \sum_{i=1}^n x_i + \frac{1}{n} n\mu_0^2\right) \\&= \frac{n}{\sigma^2} \left(2\bar{x}^2 - \bar{x}^2 - 2\mu_0\bar{x} + \mu_0^2\right) \\&= \frac{n}{\sigma^2} \left(\bar{x} - 2\mu_0\bar{x} + \mu_0^2\right) = n\left(\frac{\bar{x} - \mu_0}{\sigma}\right)^2\end{aligned}$$

Likelihood Ratio Test (LRT)

And by the CLT, we know that $\sqrt{n}\left(\frac{\bar{x}-\mu_0}{\sigma}\right) \sim N(0, 1)$.

Since, by definition,

$$\sum_{i=1}^n Z_i^2 \sim \chi_n^2$$

It follows that

$$2\ln(\Lambda(x_1, \dots, x_n)) \sim \chi_1^2$$

And therefore H_0 is rejected if $n\left(\frac{\bar{x}-\mu_0}{\sigma}\right)^2 > \chi_{1,1-\alpha}^2$.

Working example 1

Suppose that we get the following sample of size 40:

```
1  [1] 11.75  9.05  6.37 11.44 10.10 15.27  9.26 15.00 10.80 10.60 12.98 10.76
2      13.69 13.99 10.19 14.10 13.79 13.36 14.69 13.72 10.77 11.12 15.66
3      12.73 10.66  6.46 10.11  9.42 16.54 20.21 11.18 15.83  9.57 11.87 10.08
4      10.68 13.52 15.39 14.93 11.63
```

The data are available here:

<https://github.com/JRigh/Likelihood-Ratio-Tests/blob/main/data.csv>

For what integer value(s) of μ_0 the LRT statistic is NOT rejected at a significance level of 5%? In other words, what are acceptable population mean for that dataset? Answer: 12 and 13 (see code below).

Working example 1 - R code

```
1 # Set seed for reproducibility
2 set.seed(2023)
3
4 # Parameters
5 mu = 12 # true mean (unknown in practice)
6 n = 40; sigma <- 3; mu0 <- seq(5, 15, by = 1); alpha <- 0.05
7
8 # Generate artificial sample from a normal distribution
9 data <- rnorm(n, mean = mu, sd = sigma) # True mean = 12
10
11 # LRT statistic
12 lrt_statistic <- n * ((mean(data) - mu0)^2 / sigma^2)
13
14 # Calculate the critical value from the chi-squared distribution
15 critical_value <- qchisq(1 - alpha, df = 1)
16
17 # Perform the Likelihood Ratio Test
18 reject_null <- lrt_statistic > critical_value
19
20 # Print results
21 results = data.frame(mu0 = mu0,
22                       lrt_statistic = lrt_statistic,
23                       decision = ifelse(lrt_statistic > critical_value, 'Yes', 'No'))
24 #      mu0 lrt_statistic decision
25 #...
26 # 6      10      22.1237569      Yes
27 # 7      11       6.7361255      Yes
28 # 8      12       0.2373829      No
29 # 9      13       2.6275293      No
30 # 10     14      13.9065645      Yes
31 # 11     15      34.0744886      Yes
```

Working example 1 - Python code

```
1 import numpy as np
2 import pandas as pd
3 from scipy.stats import chi2
4
5 # import the data
6 data = pd.read_csv("path/data.csv")
7
8 # Calculate LRT statistic for each mu0
9 lrt_statistic = n * ((np.mean(data['x']) - mu0)**2 / sigma**2)
10
11 # Calculate the critical value from the chi-squared distribution
12 critical_value = chi2.ppf(1 - alpha, df=1)
13
14 # Perform the Likelihood Ratio Test and make decisions
15 decisions = np.where(lrt_statistic > critical_value, "Yes", "No")
16
17 # Create a results data array
18 results = np.column_stack((mu0, lrt_statistic, decisions))
19 results
20
21 # ['10', '22.136480277777775', 'Yes'],
22 # ['11', '6.743146944444444', 'Yes'],
23 # ['12', '0.23870249999999998', 'No'],
24 # ['13', '2.623146944444445', 'No'],
25 # ['14', '13.896480277777778', 'Yes'],
```

Mixture models: introduction

Let Y be a random variable and y be any observed values of this random variable. Then Y obeys a finite mixture distribution if its density can be written as $f(y) = \lambda_1 f_1(y) + \dots + \lambda_k f_k(y) = \sum_{j=1}^k \lambda_j f_j(y)$ provided that $\lambda_j > 0$ and $\sum_{j=1}^k \lambda_j = 1$. The weights λ_j are called the *mixing proportions* and $f_j(y)$ are called the *component densities*. Further, a k -component finite mixture model has the form:

$$f(y \mid \Psi) = \sum_{j=1}^k \lambda_j f_j(y \mid \theta_j)$$

In the case of a two-component mixture, that is with $k = 2$, we have the mixture parameter vector is $\Psi = (\lambda_1, \lambda_2, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2)$; the number of components is k ; the component density parameters are $\theta_1 = (\mu_1, \sigma_1^2)$ and $\theta_2 = (\mu_2, \sigma_2^2)$; the mixing proportions are λ_1 and $\lambda_2 = (1 - \lambda_1)$.

Detection of the number of components k - (1/2)

Among the many methods for determining the number of components in a mixture model, the Likelihood Ratio Test (LRT) is one commonly used in practice. A good explanation is given in McLachlan and Peel (2000). Roughly, the procedure proposes to test sequentially for a null hypothesis of a model with the smallest number of component k in the mixture against an alternative hypothesis of $k + 1$ components, that is:

$$H_0 : k = k_0 \quad \text{against} \quad H_1 : k = k_0 + 1$$

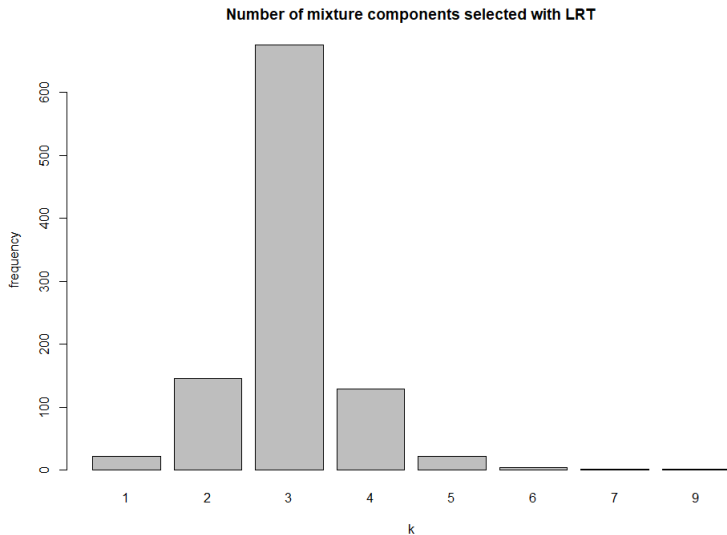
The test statistic considered is:

$$-2\log(\Lambda) = 2(\log L(\hat{\Psi}_1) - \log L(\hat{\Psi}_0)) ,$$

Detection of the number of components k - (2/2)

with $\hat{\Psi}_1$ and $\hat{\Psi}_0$ being the maximum likelihood estimates for the model with the largest respectively the smallest number of components. A small value for γ or similarly a large value for $-2\log(\gamma)$ will indicate strong evidence against the null hypothesis. The approximate distribution of the statistic under H_0 is not clearly identified, that is why we often resort to bootstrap in practice to take a decision with regards to the null hypothesis. A parametric bootstrap for this statistic is implemented in the mixtools package. This procedure is practically useful because it can be carried prior an EM or a bayesian analysis to have a guess about k , unlike information criteria, introduced in the next section, that are used conditionally to a preliminary parameter estimation.

k as selected by the LRT



Detection of k - R code

```
1 library(mixtools)
2
3 set.seed(2023)
4 data2 = rnormmix(n = 100, lambda = c(0.2, 0.3, 0.5), mu = c(1,5,8), sigma = c
      (1,1,1))
5
6 ## we will first run it 1,000 times and record each time the number of
      components selected by the lrt
7 ## hoping that one value for 'k' clearly stands out.
8 set.seed(2)
9 count.k = numeric(1000)
10 for(i in 1:1000) {
11   count.k[i] <- length(boot.comp(y=data2, max.comp=10, B=5,
12     sig=0.05, mix.type=c("normalmix"))$p.values)
13 }
14 count.k
15
16 sum(count.k==1)/1000 # percentage that LRT detect 3 components
17 sum(count.k==2)/1000 # percentage that LRT detect 3 components
18 sum(count.k==3)/1000 # percentage that LRT detect 3 components
19 sum(count.k==4)/1000 # percentage that LRT detect 4 components
20 sum(count.k==5)/1000 # percentage that LRT detect 5 components
21
22 # visualize the results
23 par(mfrow=c(1,1))
24 barplot(table(count.k), col="grey", xlab="k",
25   ylab="frequency", main="Number of mixture components selected with LRT")
26
27 # from the barplot the LRT detect the presence of k=3 subpopulations in the
      mixture.
```

References

Bijma, F., Jonker M., Van der Vaart, A., An Introduction to Mathematical Statistics. Amsterdam University Press., 2016

R.V.Hogg and E.A.Tanis: Probability and Statistical Inference, Sixth Edition, Prentice Hall, Upper Saddle River, N.J., 2001.

McLachlan, G. and Peel, D., Finite mixture models. 0471006262, John Wiley & Sons., 2000

The R Project for Statistical Computing:
<https://www.r-project.org/>

Python:
<https://www.python.org/>

course notes